

Standard Operating Procedure

Creating an Evidence-Led Monthly AI Threat Intelligence Briefing Slide in Manus

Version: 1.0

Purpose: Produce one executive, PowerPoint-compatible AI Threat Intelligence Briefing slide for a bank Risk Committee.

Audience: Any Manus user who can create a Manus account, start a task, paste a prompt, and download a presentation. No programming, threat-intelligence, or slide-production experience is assumed.

Operating principle: The slide is the end product, not the research record. Research sources, facts, limitations and 2LoD analysis must be retained in task files, while the slide contains only decision-useful conclusions and no citations or source names.

1. Outcome, Scope and Responsibilities

Use this SOP to produce a single 16:9 executive slide covering a defined full calendar month. It can also produce one connected slide covering two adjacent full months when that exception is expressly requested. The July-August 2026 example in this SOP is a two-month exception, not a rolling period.

The process converts technical developments into risk management language in this fixed sequence:

Technical intelligence finding	Required executive translation
Technical vector	What changed in a technology, identity, model, provider or dependency.
Operational vulnerability	Which banking control, business process or service could fail.
Systemic or concentration risk	How common suppliers, shared infrastructure, data, identities or platforms could propagate the effect.
Geopolitical or macro-financial threat	How actor intent or cross-border dependency could affect confidence, continuity or financial stability.

Financial or regulatory impact

How operational loss, customer harm, service interruption, conduct failure, enforcement or supervisory action could arise.

The person requesting the slide owns the coverage period, intended audience and any bank-specific business context. Manus performs research, evidence validation, drafting and presentation creation. The requestor conducts the final management review before external or committee circulation.

2. Prerequisites and Working Rules

Create a Manus account and open a new task. The procedure uses only public information during research, so no customer information, internal incidents, credentials, source-code extracts or confidential board materials should be added to the task unless the task has separately been approved for that use.

Install two custom skills before starting production:

1. `monthly-ai-threat-intelligence-briefing` controls the evidence standard, six research streams, executive risk translation and mandatory slide layout.
2. `risk-committee-writing-controls` prevents rhetorical risk language and enforces a deterministic language check.

Use a new task for every reporting period. Preserve an internal evidence register within that task, because it enables challenge, auditability and later updates. Do not reuse a prior slide as evidence. Reuse its design only after rebuilding the evidence record for the new reporting period.

Item	Required standard
Coverage basis	Full calendar month or an expressly requested fixed multi-month period. Never use a rolling 30-day window.
Evidence threshold	At least three directly read primary-source publications underpin the slide.
Source hierarchy	Financial supervisors and regulators; national cyber agencies; impacted-entity disclosures; law-enforcement or sanctions notices; then high-quality threat intelligence for corroboration.

Selection threshold	Include only active exploitation, material incident/impact, high-dependency common technology, banking relevance, regulatory deadline/enforcement effect, or credible cross-border consequence.
Slide density	Prioritise fewer, substantiated developments over smaller text or unsupported claims.
Output	Exactly one PowerPoint-compatible slide, except where the requestor expressly asks for a different deliverable.

3. Build and Install the Custom Skills

3.1 No-code method - recommended for a first-time user

Create a new Manus task titled **Create Risk Committee briefing skills**. Paste the two creation prompts below, one at a time. Manus will create an installable skill package and return it as an attachment card. Use **Add to My Skills** on each card. After both cards show as installed, open a new task for the monthly briefing.

The first prompt creates the writing-control skill. It must contain a small deterministic checker so language controls can be tested before and after slide creation.

Plain Text

```
Create an installable Manus skill named risk-committee-writing-controls.

Use the Skill Creator workflow. Create a concise required SKILL.md and a script

Apply these controls to every user-visible sentence, including headings, labels
- Do not use the adjectives serious, critical, extreme, or verified, except tha
- Do not use an em dash. Use a spaced hyphen: " - ".
- Separate observed activity, assessed exposure and plausible scenario.
- Prefer material exposure, active exploitation, unauthenticated remote access,
- Do not deliver the final output until the checker passes on the draft content

The checker must exit non-zero when it finds a prohibited word or an em dash. D
```

The second prompt creates the monthly briefing skill. It contains the research workflow, its detailed research protocol and the single-slide blueprint.

Plain Text

Create an installable Manus skill named monthly-ai-threat-intelligence-briefing

Use the Skill Creator workflow. Create a concise required SKILL.md, a reference

The SKILL.md must instruct Manus to:

1. Determine the previous full calendar month unless the user expressly supplies
2. Run six research streams exactly as set out in research_protocol.md.
3. Read every selected primary source directly, not just search snippets.
4. Select no more than five developments and use at least three primary-source
5. Exclude generic AI releases, laboratory-only claims, vendor marketing, unsupported
6. Translate every selected item from technical vector to operational vulnerability
7. Make a definite 2LoD risk-posture conclusion without asserting an internal bias
8. Use the exact six-part single-slide layout in slide_blueprint.md.
9. Produce exactly one PowerPoint-compatible 16:9 slide and no on-slide citations
10. Use the risk-committee-writing-controls skill and its checker before and after

Deliver the completed installable skill package.

3.2 Manual skill contents - use when you need to inspect or amend the design

A Manus skill is a small folder with one required file, SKILL.md, and optional bundled files. The folder layout below is sufficient for this use case. The folder and file names must remain lower-case and use hyphens or underscores as shown.

Plain Text

```
risk-committee-writing-controls/
├── SKILL.md
└── scripts/
    └── check_language_controls.py

monthly-ai-threat-intelligence-briefing/
├── SKILL.md
├── references/
│   └── research_protocol.md
└── templates/
    └── slide_blueprint.md
```

If a Manus task is asked to build these skills, it should first initialise the skill package using its standard creation workflow, then replace the generated placeholders with the content in the following sections. Remove unused example files. The task should validate each package before returning it for installation.

3.3 risk-committee-writing-controls/SKILL.md

Markdown

```
---
name: risk-committee-writing-controls
description: Apply mandatory language controls to Risk Committee reports, presentations, etc.
---

# Risk Committee Writing Controls

Apply these controls to every user-visible sentence in reports, presentations, etc.

| Control | Required treatment |
|---|---|
| Restricted vocabulary | Do not use `serious`, `critical`, `extreme`, or `verified`. |
| Dash character | Do not use an em dash. Use a spaced hyphen: ` - `. |
| Evidence discipline | State observed activity, assessed exposure and plausible outcomes. |
| Risk language | Describe materiality, exploitation status, privilege, reachability, etc.

Use source-supported, observed, assessed, not established by available evidence.

Before delivery, run `python scripts/check_language_controls.py <text-file>` on
```

3.4 risk-committee-writing-controls/scripts/check_language_controls.py

Python

```
#!/usr/bin/env python3
import re
import sys
from pathlib import Path

if len(sys.argv) != 2:
    print("Usage: python check_language_controls.py <text-file>")
    raise SystemExit(2)

path = Path(sys.argv[1])
text = path.read_text(encoding="utf-8")
scan_text = text.replace("Critical Vulnerability Tracking", "")
issues = []

for match in re.finditer(r"\b(serious|critical|extreme|verified)\b", scan_text):
    line = scan_text.count("\n", 0, match.start()) + 1
    issues.append(f"Line {line}: restricted term '{match.group(0)}'")
```

```

for match in re.finditer("-", scan_text):
    line = scan_text.count("\n", 0, match.start()) + 1
    issues.append(f"Line {line}: em dash found; replace with ' - ")

if issues:
    print("FAIL - language controls breached")
    print("\n".join(issues))
    raise SystemExit(1)

print("PASS - language controls satisfied")

```

3.5 monthly-ai-threat-intelligence-briefing/SKILL.md

Markdown

```

---
name: monthly-ai-threat-intelligence-briefing
description: Produce an evidence-led monthly AI Threat Intelligence Briefing for
---

# Monthly AI Threat Intelligence Briefing

Produce one executive 16:9 PowerPoint slide for a Group CRO, Risk Committee and

> Technical vector → operational vulnerability → systemic/concentration risk →

Read `references/research_protocol.md` before research and `templates/slide_blue

## Workflow

1. Determine the coverage period. Use the prior full calendar month unless the
2. Run the six research streams in the research protocol.
3. Validate centrally. Read every selected primary source and retain the intern
4. Select material developments only. Use at least three primary-source publica
5. Form the 2LoD judgement. Challenge first-line assumptions on external-asset
6. Prepare the slide using the exact six sections and the five-block risk chain
7. Generate exactly one PowerPoint-compatible slide in 16:9 format.
8. Run the writing-controls checker before content creation and against complet

```

3.6 monthly-ai-threat-intelligence-briefing/references/research_protocol.md

Markdown

```

# Research Protocol

## Coverage and source standard

```

Search only for developments published, disclosed, observed or materially updated.
Every shortlisted item needs: event date; publication date; named primary source.

Six mandatory research streams

1. Supervisory and policy developments: material AI, cyber, operational-resilience.
2. AI-native threat actor activity: substantiated campaigns involving AI-systems.
3. AI systems, applications and supply chain: exploitation, vulnerabilities, controls.
4. Material incidents and financial-sector impact: substantiated incidents affecting financial systems.
5. Systemic concentration and macro-financial exposure: provider concentration, dependencies.
6. Defensive efficacy and remediation intelligence: demonstrably effective controls.

Standard research prompt

> Research only material AI-related cyber-risk developments for ****[coverage period]**.

Central validation

Select no more than five developments. At least three primary-source publications.

Risk chain

Create a five-block risk chain only where a single pathway is evidence-supported.

****Actor/Intent → AI Enabler or Initial Access → Exploitation/Pivot → Operational Impact**

If no evidence-supported linear chain exists, state: `No evidence-supported end state`.

3.7 monthly-ai-threat-intelligence-briefing/templates/slide_blueprint.md

Markdown

Single-Slide Blueprint

Design system

Element	Standard
Canvas	16:9, 1280 × 720px.
Background	Solid white `#FFFFFF`.
Body text	`#111111`, Inter or equivalent sans serif, minimum 16pt.
Accent	Emerald green `#00875A` only for title, executive-summary band, section headers.
Structure	Thin `#E5E7EB` rules. Do not use gradients, shadows, dense cards, or complex styling.

Required layout

1. **Title:** `AI THREAT INTELLIGENCE BRIEFING - [PERIOD]`, upper left in emerald
2. **Executive summary band:** A full-width emerald band below the title with one line of text
3. **Left column (31%):** `EVENTS OF THE MONTH` with 2-3 dated developments. Each event in a callout box
4. **Centre column (30%):** `PRIMARY AI RISK VECTOR` as a vertical five-block chain
5. **Right column (39%):** `SYSTEMIC EXPOSURE & MATERIAL INCIDENTS`. Put concentration on the most significant incident
6. **Bottom strip:** `REMEDIAL ACTION GUIDANCE`, divided into three equal action cells

Content limits

Area	Maximum
Executive summary	35 words.
Events	3 events; 25 words each.
Risk chain	5 blocks; 12 words per block.
Systemic/incidents	4 callouts; 25 words each.
Remediation	3 directives; 18 words each.
Citations/metadata	None on the slide. Keep sources only in the internal evidence

4. Start a Monthly Briefing Task

Open a fresh Manus task after the two skills have been installed. Invoke both skills in the first line. State the period, the target audience, the required number of slides and whether the period is an exception to the normal single-month cycle.

For a standard monthly briefing, use this prompt.

Plain Text

```
/monthly-ai-threat-intelligence-briefing /risk-committee-writing-controls
```

Create exactly one evidence-led, PowerPoint-compatible AI Threat Intelligence Briefing

For the required July-August connected slide, use this prompt.

Plain Text

```
/monthly-ai-threat-intelligence-briefing /risk-committee-writing-controls
```

Create exactly one connected, evidence-led, PowerPoint-compatible AI Threat Intelligence

The response should acknowledge the task and then begin evidence collection. If the scope, coverage dates or desired output format are ambiguous, clarify them before drafting. Do not ask for technical data when public evidence is sufficient.

5. Research and Evidence Validation

5.1 Run the six streams

The six research streams prevent the slide becoming a collection of vendor claims or generic AI news. They must be handled separately because they represent different risk pathways: supervisory exposure, external threat activity, technical/supply-chain exposure, financial-sector impact, systemic concentration and defensive control effectiveness.

Manus should research all six streams, then centralise the output. The research prompt in Section 3.6 is the unit-level prompt. The central task must require each stream to return candidate facts, direct URLs, banking relevance, risk translation and evidence limitations.

5.2 Create the internal evidence register

Use a task file, spreadsheet or structured note. Keep it internal. The register below is the minimum schema.

ID	Development	Event date	Publication date	Primary source and URL	Source-supported fact
V1	[Short name]	[DD Mon YYYY]	[DD Mon YYYY]	[Authority and direct URL]	[One factual sentence]

Follow this decision logic:

Candidate status	Required treatment
Primary source has been directly read and supports an in-period material development	Eligible for selection.
Primary source exists but describes only a future proposal, a generic product release or a non-banking scenario	Context only or exclude.
Media, vendor report or social post makes a claim without an accessible primary source	Exclude from the slide. Use only as a discovery lead.
Public data cannot show whether any bank was affected	State the evidence limitation. Do not infer internal impact.

A threat pathway is technically plausible but no public event substantiates it

Record it as an assessed exposure, not as an incident.

5.3 Read the sources, do not rely on result snippets

Search results are leads, not evidence. Open each selected direct source and confirm its publication date, the affected entity or sector, the event date, what actually happened, and the scope of the source's own claims. Save key extracts or a concise record to the internal evidence register while researching.

Do not turn probability, vendor marketing or a laboratory demonstration into a material financial-sector incident. A laboratory event can still be included when it directly proves a control pathway relevant to banks, but its limitation must be stated: for example, internal evaluation, non-production model, or absence of a disclosed customer or bank impact.

5.4 Select no more than five developments

The final slide should usually contain two or three dated events, one connected risk chain, two to four systemic callouts, and three management actions. The same selected source can inform more than one section, but do not duplicate its wording.

The evidence register must show at least three primary-source publications across the selected narrative. A source can underpin remediation guidance even when it is outside the coverage period only if it remains a valid active supervisory or control expectation. It must not be presented as a reporting-period event in that case.

6. Produce the 2LoD Risk Judgement

Write one conclusion of 25-35 words for the emerald summary band. The conclusion should identify the changed risk posture and the management response. It should not recount every event.

Use the following method:

1. **Observed activity:** State what the source establishes. Example: "A model evaluation bypassed isolation controls and accessed a third-party system."
2. **Assessed exposure:** State the bank-relevant control weakness or concentration pathway. Example: "AI agents with broad tool access, persistent privileges or unrestricted egress could extend a local control gap across shared data and cloud services."

3. **Plausible impact:** State the business outcome conditionally. Example: “This could create service disruption, sensitive-data exposure, conduct-control failure and increased supervisory engagement.”
4. **RAF position:** Convert the judgement into measurable risk tolerances, evidence requirements or action owners. Example: “Set measurable tolerance for agent authority, provider substitutability, third-party reachability, content provenance and recoverability.”

The primary risk vector must be linear and evidence-supported. If the sources do not support a full pathway, do not force one. Use the evidence-gap statement from the research protocol and make the centre column a structural-risk explanation instead.

7. Prepare the Slide Copy Before Presentation Creation

Create one short content document before generating the slide. The content document has no citations because the evidence record retains them. Run the writing-control checker at this stage.

Use this completion template.

Markdown

```
# AI THREAT INTELLIGENCE BRIEFING - [PERIOD]

## Executive summary
[25-35 word 2LoD conclusion]

## Events of the month
**[DD MON] | [EVENT LABEL]**
[One business implication sentence - 25 words maximum]

**[DD MON] | [EVENT LABEL]**
[One business implication sentence - 25 words maximum]

## Primary AI risk vector
**Actor / Intent:** [12 words maximum]
**AI Enabler or Initial Access:** [12 words maximum]
**Exploitation / Pivot:** [12 words maximum]
**Operational Compromise:** [12 words maximum]
**Financial & Regulatory Impact:** [12 words maximum]

## Systemic exposure & material incidents
**[CALLOUT LABEL]**
[25 words maximum]
```

****[CALLOUT LABEL]****
[25 words maximum]

****[EVIDENCE LIMITATION]****
[25 words maximum]

****2LoD RAF VIEW****
[25 words maximum]

Remediation guidance

****[2-4 WORD ACTION]****
[18 words maximum]

****[2-4 WORD ACTION]****
[18 words maximum]

****[2-4 WORD ACTION]****
[18 words maximum]

The controls below must be applied before the slide is generated.

Control	Check
Language	The checker passes. Remove the restricted words and em dash character; use a spaced hyphen where necessary.
Dates	Every event date is in the requested reporting period.
Evidence	Each factual claim can be traced to the internal register.
Certainty	Observed facts, assessment and plausible effect remain distinguishable.
Scope	No claim of an internal bank failure, customer loss or regulatory breach without bank-specific evidence.
Density	Word limits are met. Do not reduce the body font below 16pt to compensate for excess content.

8. Create the PowerPoint-Compatible Slide in Manus

Ask Manus to create a one-slide, 16:9 PowerPoint-compatible presentation, not a flat image. Give it the final content document and instruct it to create the one-slide outline before laying out the content. The single slide must follow the visual blueprint exactly. Use this presentation instruction if the task needs an explicit prompt.

Plain Text

Create the final presentation as exactly one editable, PowerPoint-compatible 16

The required layout and spacing are summarised below.

Slide area	Location and content	Design requirement
Title	Upper left. AI THREAT INTELLIGENCE BRIEFING - [PERIOD] .	Emerald text; period appears here only.
Executive summary band	Full width directly below title.	Emerald fill; white one-sentence conclusion.
Left column	31% of main body. Two or three dated events.	Emerald section header; dark text; thin dividers.
Centre column	30% of main body. Five vertically connected blocks.	Near-black blocks with emerald boundaries; green end-state emphasis.
Right column	39% of main body. Concentration/regulation first, then material incidents and explicit limitation.	Emerald labels; dark text; thin dividers.
Bottom strip	Full width; three equal action cells.	Emerald heading; three short directives; no card styling.

The presentation output should be inspected visually in Manus before delivery. Confirm that no text runs outside a cell or overlaps another element; headings fit on one line; the centre chain is vertically connected; and all body text is legible without zooming. If the slide is too dense, remove the lowest-value callout rather than shrinking text or increasing the number of slides.

9. Final Quality Gate and Delivery

Run the language checker again on the completed slide text or its presentation markup. Then complete the following quality gate before presenting the final slide.

Quality gate	Pass condition
Coverage dates	The title and dated events match the defined full calendar period.
Evidence standard	At least three directly read primary-source publications underpin the selected narrative.
Incident discipline	No named incident, affected bank, loss, actor attribution or control failure appears without source support.
Evidence limitation	An explicit limitation appears if no production financial-sector incident is substantively established.
2LoD posture	The slide includes one RAF-focused challenge to first-line assumptions and three evidence-oriented actions.
Risk chain	The five-block chain describes one source-supported pathway, or clearly states that no end-to-end chain is supported.
Writing controls	The pre-generation and post-generation language checks pass.
Slide controls	No citations, source names, logos, footers, page numbers, audience labels, classification tags, watermarks, gradients, imagery or shadows.
Output	Exactly one 16:9 PowerPoint-compatible slide is available for download.

Deliver only the presentation output unless the requestor asks for supporting evidence. Keep the evidence register in the task files for internal challenge and reproducibility.

10. Exception Handling and Troubleshooting

Situation	Required response
-----------	-------------------

The user asks for two months on one slide	Treat it as an explicit fixed-period exception. Retain the same source and quality standard. Use both months only in the title and maintain one connected narrative.
Fewer than three primary sources are available	Do not fabricate a third source. Present a lower-density slide, include the evidence limitation, and retain only substantiated structural or supervisory developments.
A selected source was published after the reporting period	Use it only when it explicitly documents an event occurring inside the period. Mark it as retrospective reporting in the internal register.
A report is persuasive but comes from a vendor	Use it only to find primary sources or as corroboration. Do not rely on it alone for an on-slide material incident.
A vulnerability or model issue has no evidence of active exploitation	Frame it as structural exposure or remediation intelligence, not as a realised incident.
The model invents a source, date, breach scale or bank impact	Exclude the claim immediately. Re-open the direct source and rebuild the related slide text.
The risk chain is not linear	Use the explicit evidence-gap statement rather than combining unrelated developments into an artificial chain.
The slide contains too much text	Reduce events to two, combine overlapping callouts, or remove a lower-value item. Do not decrease body text below the readability floor.
The final slide has citations or sources	Remove all references from the slide and retain them only in the internal evidence register.
The language checker fails	Replace each flagged word or dash, re-read the sentence for proportionality, and run the checker again.
The slide renders as an image rather than a presentation	Regenerate it as a PowerPoint-compatible presentation and confirm that the downloadable deliverable supports slide export.

11. Monthly Operating Checklist

Use this one-page sequence for each run after the two skills have been installed.

1. Open a new Manus task and invoke both custom skills.
2. State the reporting period, target audience and “exactly one PowerPoint-compatible slide” .
3. Run all six research streams.
4. Capture candidates in the evidence register.
5. Open and read every source proposed for selection.
6. Select no more than five developments and retain at least three primary-source publications.
7. Create the five-block chain only when one pathway is evidence-supported.
8. Write a 2LoD RAF conclusion and three evidence-oriented actions.
9. Prepare slide copy using the six-section template and word limits.
10. Run the language checker on the copy.
11. Generate the one-slide 16:9 PowerPoint-compatible presentation.
12. Inspect the presentation visually and run the language checker again on completed slide text.
13. Complete the final quality gate and deliver the slide.

Completion standard: The delivered slide enables a Risk Committee to distinguish observed AI-related developments, assessed bank exposure, plausible financial/regulatory pathways, and specific evidence requirements for management action.